

CS331: Java Threads

<http://jatinga.iitg.ac.in/~asahu/cs331/>

A Sahu

Dept of Computer Science & Engineering

IIT Guwahati

Outline

- Thread Methodology
- Eight Rules for Designing Multithreaded Apps
- Java thread
 - Creation and lifecycle
- Examples
 - VectorSum
- Synchronized
 - Function, block, objects
- **Source Code Available@**
<http://jatinga.iitg.ac.in/~asahu/cs331/>

Multicore Difficulties

- Multiprocessors are likely to be cost/power effective solutions
 - Because it share lots of resources
 - *Personal room is costlier than dormitory*
 - Sharing resource arise many other problems
 - Critical Sections
 - Lock and Barrier Design
 - Coherence
 - Shared data at all placed should be same
 - Consistency
 - Order should be similar to serial (ROB)
 - One processor Interference others
 - Share efficiently using some policy

Threading Methodology

Threading Methodology

- **Does not recommend going straight to concurrency!**
- First produce a tested single-threaded program
 - Use reqs./ design/ implement /test/ tune/ maintenance steps
- Then to create a concurrent system from the former, do
 1. Analysis: Find computations that are independent of each other
 1. AND take up a large amount of serial execution time (80/20 rule)
 2. Design and Implement: straightforward Test for Correctness: Verify that concurrent code produces correct output
 3. Tune for performance: once correct, find ways to speed up

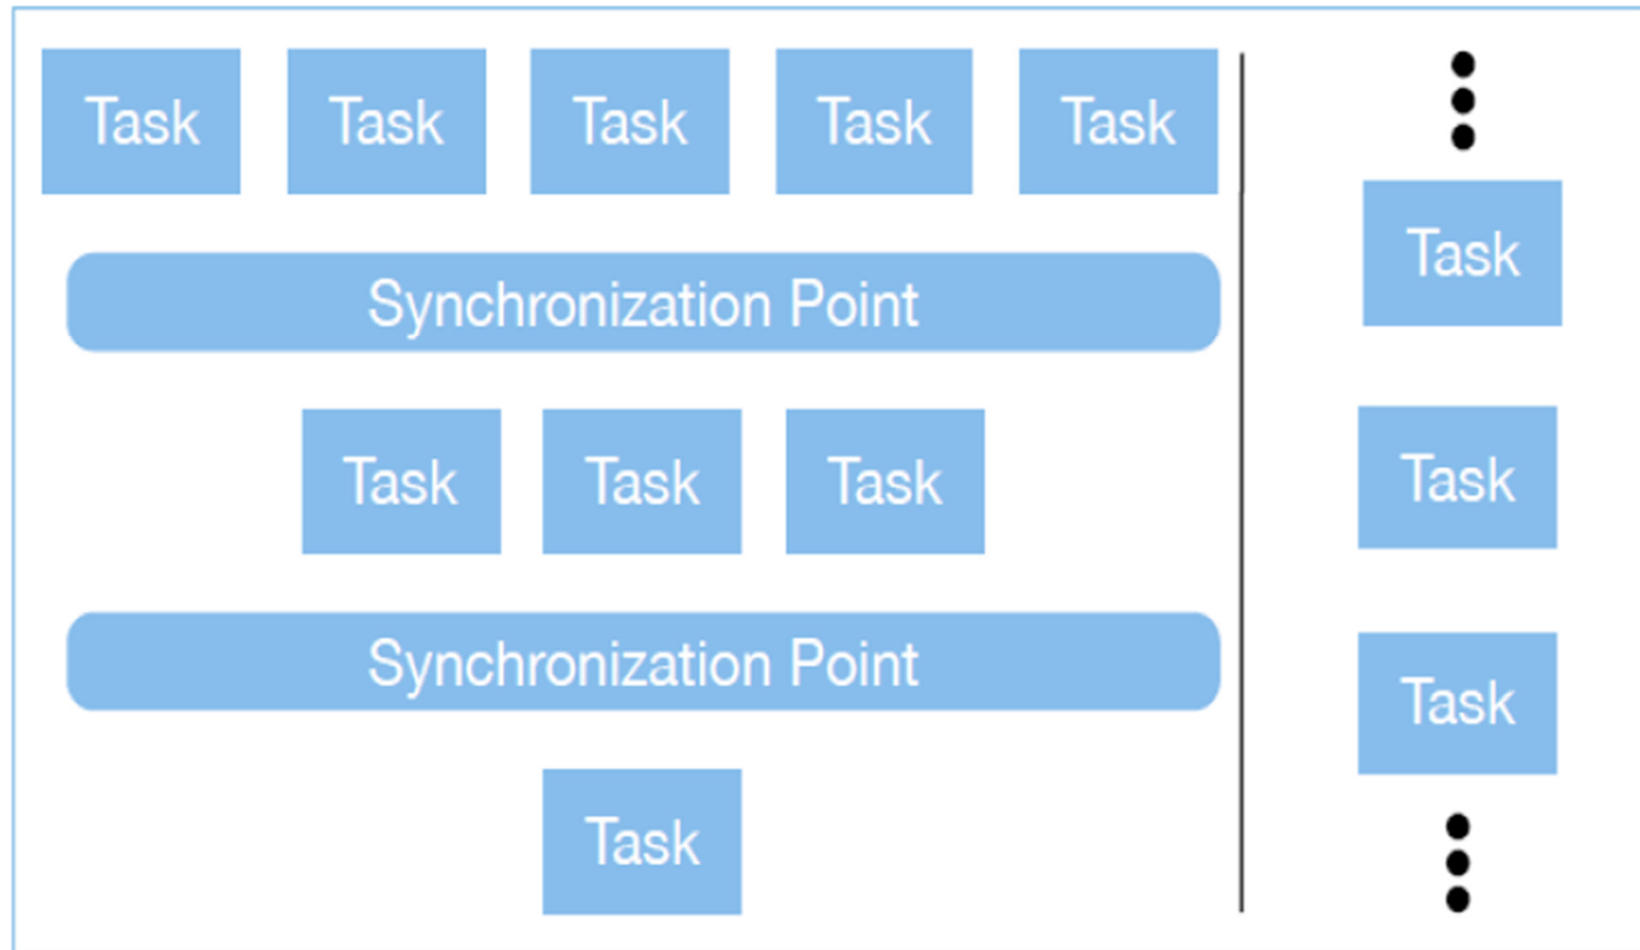
Performance Tuning

- Tuning threaded code typically involves
 - identifying sources of contention on locks (synchronization)
 - identifying work imbalances across threads
 - reducing overhead
- Testing and Tuning
 - Whenever you tune a threaded program, you must test it again for correctness
- Going back further
 - if you are unable to tune system performance,
 - you may have to re-design and re-implement

Design Models

- Two primary design models for concurrent algorithms
- Task Decomposition
 - identify tasks (computations) that can occur in any order
 - assign such tasks to threads and run concurrently
- Data Decomposition
 - program has large data structures where individual data elements can largely be calculated independently
 - data decomposition implies task decomposition in these cases

Task Decomposition



concurrent system ← sequential system

Eight rules of Designing multithreaded APPS

1/8 Rules: Designing MT APPS

- **Identify Truly Independent Computations**
- If you can't identify (in a single threaded application) computations that can be done in parallel, you're out of luck
- Some situations that indeed can't be made parallel

2/8 Rules: Designing MT APPS

- **Implement Concurrency at the Highest Level Possible**
- When discussing “What’s Not Parallel” a common refrain was “you can’t make this parallel,
 - So see if its part of a larger computation that CAN be made parallel”
- This is such good advice, it was promoted to being a guideline!
 - Two approaches: bottom up, top down

2/8 Rules: Bottom UP

- One methodology says to create a concurrent program
 - Start with a tuned, single-threaded program and
 - Use a profiler to find out where it spends most of its time
- In the bottom-up approach, you start at those “hot spots” and work up; typically, a hotspot will be a loop of some sort
 - See if you can thread the loop
 - If not, move up the call chain, looking for the next loop and see if it can be made parallel...
 - If so, still look up the call chain for other opportunities, first.
- Why? Granularity! You want coarse-grained tasks for your thread

2/8 Rules: Top Down

- With knowledge of the location of the hot spot
- Start by looking at the whole application and see if there are parallelization opportunities on the large-scale structure that contains the hot spot
 - if so, you've probably found a nice coarse-grained task to assign to your threads
 - If not, move lower in the code towards the hot spot, looking for the first opportunity to make the code concurrent

3/8 Rules: Designing MT APPS

- **Plan Early for Scalability**
- The number of cores will keep increasing
- You should design your system to take advantage of more cores as they become available
 - Make the number of cores an input variable and design from there
- In particular, designing systems via data decomposition techniques will provide more scalable systems
 - humans are always finding more data to process!
- More data, more tasks; if more cores arrive, you're ready

4/8 Rules: Designing MT APPS

- **Make use of Thread-Safe Libraries Wherever Possible**
- First, software reuse!
 - Don't fall prey to Not Invented Here Syndrome
 - if code already exists to do what you need, use it!
- Second, more libraries are becoming multithread aware
 - That is, they are being built to perform operations concurrently
- Third, if you make use of libraries, ensure they are thread-safe; if not, you'll need to synchronize calls to the library
 - Global variables hiding in the library may prevent even this, if the code is not reentrant ; if so, you may need to abandon it

5/8 Rules: Designing MT APPS

- **Use the Right Threading Model**
- Avoid the use of explicit threads if you can get away with it
- They are hard to get right
- Look at libraries that abstract away the need for explicit threads
 - OpenMP, Cilk and Intel Threading Building Blocks
 - Scala's agent model, Go's go routines and Clojure's concurrency primitives
- All of these models hide explicit threads from the programmer Right Threading Model

6/8 Rules: Designing MT APPS

- **Never Assume a Particular Order of Execution**
- With multiple threads, as we've seen, the scheduling of atomic statements is nondeterministic
- If you care about the ordering of one thread's execution with respect to another, you have to impose synchronization
- But, to **get the best performance**, you want to **avoid synchronization** as much as possible
- In particular, you want high granularity tasks that don't require synchronization
 - This allows your cores to run as fast as possible on each task they're given

7/8 Rules: Designing MT APPS

- **Use Thread-Local Storage Whenever Possible or Associate Locks with specific data**
- Related to Rule 6; the more your threads can use thread-local storage, the less you will need synchronization
- **Otherwise, associate a single lock with a single data item** —
 - in which a data item might be a huge data structure
- This makes it easier for the developer to understand the system;
 - “if I need to update data item A, then I need to acquire lock A first”

8/8 Rules: Designing MT APPS

- **Dare to Change the Algorithm for a Better Chance of Concurrency**
- Sometimes a tuned, single-threaded program makes use of an algorithm which is not amenable to parallelization
- They might have picked that algorithm for performance reasons
 - Strassen's Algorithm $O(n^{2.81})$ vs. the triple-nested loop algorithm to perform matrix multiplication $O(n^3)$
- Change the algorithm used by the single-threaded program to see if you can then make that new algorithm concurrent
 - BUT: when measuring speedup, compare to the original!!

Java Threads

Creating Threads

- There are two ways to create our own Thread object
 1. Subclassing the Thread class and instantiating a new object of that class
 2. Implementing the Runnable interface
- In both cases the `run()` method should be implemented

Extending Thread

```
public class ThreadExample extends Thread {  
    public void run () {  
        for (int i = 1; i <= 100; i++) {  
            System.out.println("Thread: " + i);  
        }  
    }  
}
```

Thread Methods

- **void start()**
 - Creates a new thread and makes it runnable
 - This method can be called only once
- **void run()**
 - The new thread begins its life inside this method
- **void stop() (deprecated)**
 - The thread is being terminated

Thread Methods

- **yield()**
 - Causes the currently executing thread object to temporarily pause and allow other threads to execute
 - Allow only threads of the same priority to run
- **sleep(int *m*)/sleep(int *m*,int *n*)**
 - The thread sleeps for *m* milliseconds, plus *n* nanoseconds

Implementing Runnable

```
public class RunnableExample implements Runnable {  
    public void run () {  
        for (int i = 1; i <= 100; i++) {  
            System.out.println ("Runnable: " + i);  
        }  
    }  
}
```

A Runnable Object

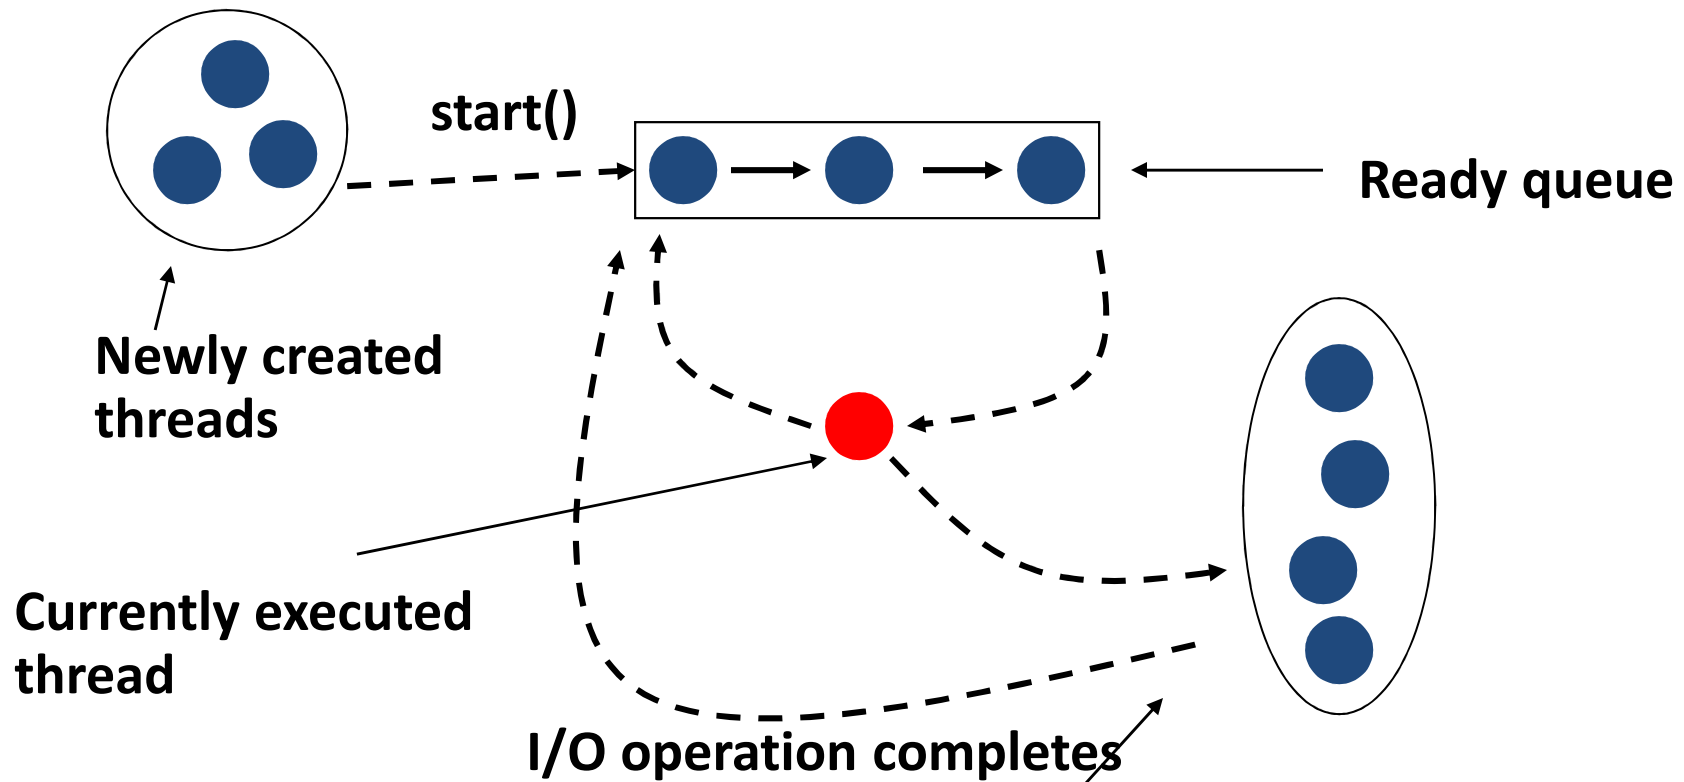
- The Thread object's `run()` method calls the Runnable object's `run()` method
- Allows threads to run inside any object, regardless of inheritance

Example - an applet that is also a thread

Starting the Threads

```
public class ThreadsStartExample {  
    public static void main (String argv[]) {  
        new ThreadExample ().start ();  
        new Thread(new RunnableExample ()).start ();  
    }  
}
```

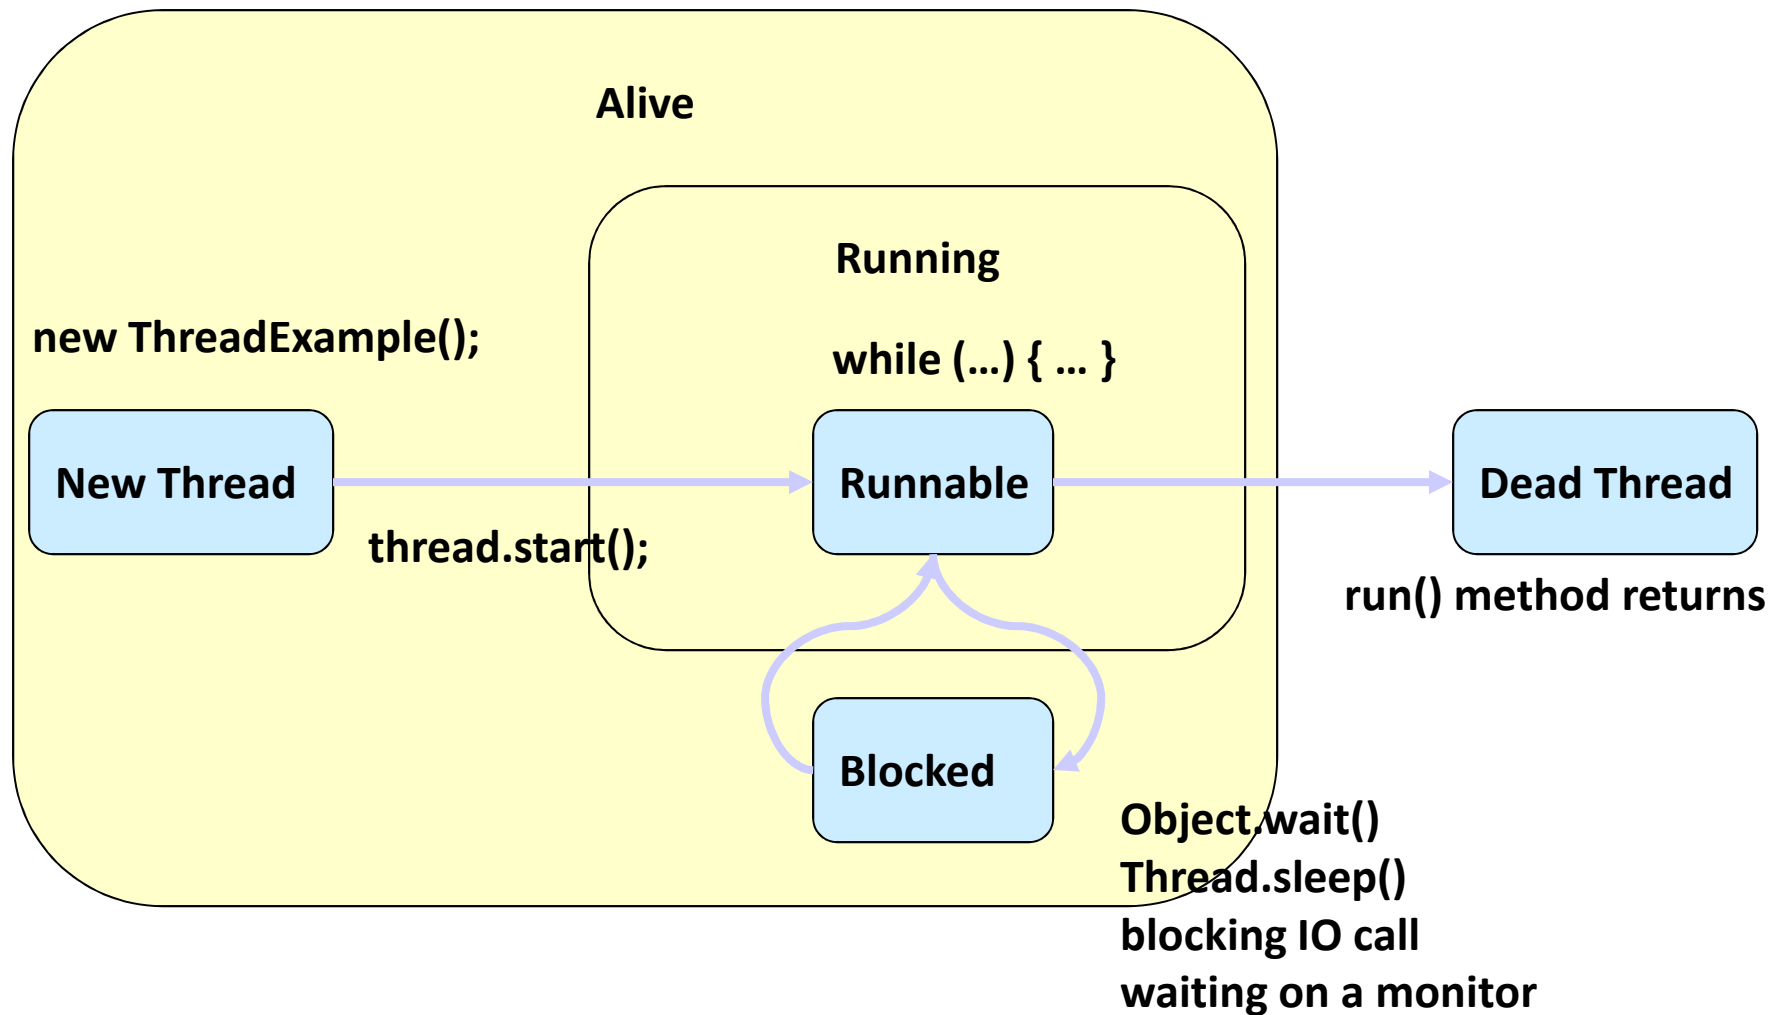
Scheduling Threads



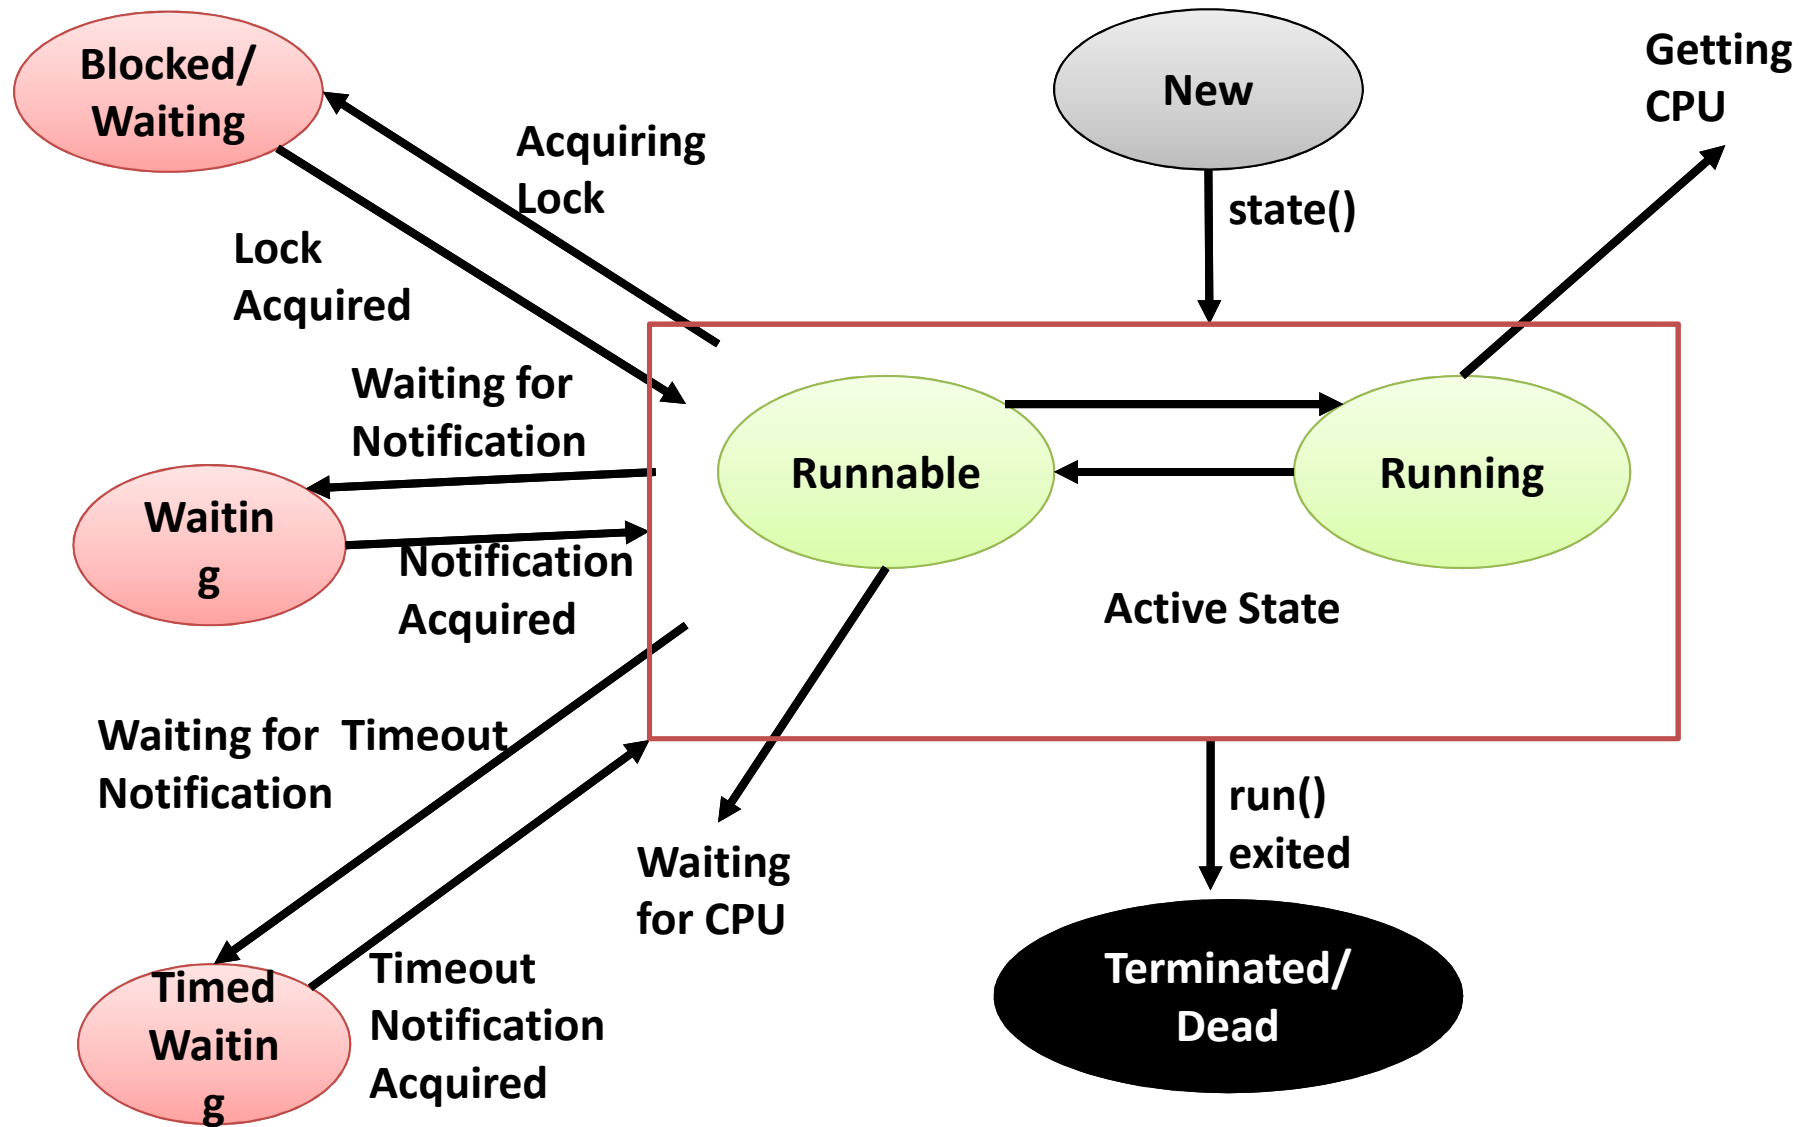
What happens when a program with a `ServerSocket` calls `accept()`?

- Waiting for I/O operation to be completed
- Waiting to be notified
- Sleeping
- Waiting to enter a synchronized section

Thread State Diagram



Java thread State: Life cycle



First Example: MainExtend.java

```
public class MainExtend extends Thread {  
    public static void main(String[] args) {  
        MainExtend thread = new MainExtend();  
        thread.start();  
        System.out.println("This code is  
                           outside of the thread");  
    }  
    public void run() {  
        System.out.println("This code is running in a thread");  
    }  
}
```

First Example: MainImplRun.java

```
public class MainImplRun implements Runnable {  
    public static void main(String[] args) {  
        MainImplRun obj = new MainImplRun();  
        Thread thread = new Thread(obj);  
        thread.start();  
        System.out.println("This code is outside of the  
thread");  
    }  
    public void run() {  
        System.out.println("This code is running in a thread");  
    }  
}
```


Example 1

```
public class PrintThread1 extends Thread {  
    String name;  
    public PrintThread1(String name) {  
        this.name = name;  
    }  
    public void run() {  
        for (int i=1; i<500 ; i++) {  
            try {  
                sleep((long)(Math.random() * 100));  
            } catch (InterruptedException ie) { }  
            System.out.print(name);  
        }  
    }  
}
```

Example 1 (cont)

```
public static void main(String args[]) {  
    PrintThread1 a = new PrintThread1("*");  
    PrintThread1 b = new PrintThread1("-");  
    PrintThread1 c = new PrintThread1("=");  
    a.start();  
    b.start();  
    c.start();  
}  
}
```

Java thread Example 2

```
class NewThread implements Runnable {  
    Thread t;  
    NewThread() {  
        t = new Thread(this, "Demo Thread");  
        System.out.println("Child thread: " + t);  
        t.start();    // Start the thread  
    }  
    public void run() {  
        for(int i = 5; i > 0; i--) {  
            System.out.println("Child Thread: " + i);  
        }  
    }  
}
```

Java final keywords

- Keyword **final** : non access modifier
- Different context final are used
 - Variable : To create constant variable
 - Method : to prevent method overriding
 - Class: to Prevent inheritance

Vector Sum Example

```
import java.util.concurrent.*;

public class VectorSum {
    private static final int MAX = 16; // Size of array
    private static final int MAX_THREAD = 4;
    private static int[] a = { 1, 5, 7, 10, 12, 14, 15, 18, 20, 22, 25, 27, 30, 64, 110, 220 };
    private static int[] sum = new int[MAX_THREAD];
    private static int part = 0;
```

static class SumArray implements Runnable {

@Override

```
    public void run() {
        // Each thread computes sum of 1/4th of array
        int thread_part = part++;
        for (int i = thread_part * (MAX / 4); i < (thread_part + 1) * (MAX / 4); i++) {
            sum[thread_part] += a[i];
        }
    }
}
```

```
}
```

Vector Sum Example

```
// Driver Code
```

```
public static void main(String[] args) throws InterruptedException {
```

```
    Thread[] threads = new Thread[MAX_THREAD];
```

```
    // Creating 4 threads
```

```
    for (int i = 0; i < MAX_THREAD; i++) {
```

```
        threads[i] = new Thread(new SumArray());
```

```
        threads[i].start();
```

```
    }
```

```
    // Joining 4 threads i.e. waiting for all 4 threads to complete
```

```
    for (int i = 0; i < MAX_THREAD; i++) { threads[i].join(); }
```

```
    // Adding sum of all 4 parts
```

```
    int total_sum = 0;
```

```
    for (int i = 0; i < MAX_THREAD; i++) { total_sum += sum[i]; }
```

```
    System.out.println("sum is " + total_sum);
```

```
}
```

```
}
```

Basic thread: Command line Args

```
import java.util.concurrent.*;
public class BasicThread {
    private static int part = 0;
    static class PrintThread implements Runnable {
        @Override
        public void run() {
            int thread_part = part++;
            char c= (char) (65+thread_part);
            for (int i = 0; i< 10; i++) { System.out.print(c); }
        }
    }
}
```

Basic thread: Command line Args

```
// Driver Code
```

```
public static void main(String[] args) throws InterruptedException {  
    for(int i=0;i<args.length;i++) System.out.println(args[i]); //parse  
    MAX_THREAD=Integer.parseInt(args[0]); //First Argument : for running  
with four thread Use command : java BasicThread 4  
    Thread[] threads = new Thread[MAX_THREAD];  
        System.out.println("Number of thread created..="+MAX_THREAD);  
    for (int i = 0; i < MAX_THREAD; i++) {  
        threads[i] = new Thread(new PrintThread());  
        threads[i].start();  
    }  
    for (int i = 0; i < MAX_THREAD; i++) { threads[i].join(); }  
}
```


Mutlithreaded Counter

```
// Java Program to demonstrate synchronization in Java
class Counter {
    private int c = 0; // Shared variable
    // Synchronized method to increment counter
    public synchronized void inc() {
        c++;
    }
    // Synchronized method to get counter value
    public synchronized int get() {
        return c;
    }
}
```

Mutlithreaded Counter

- Contains a private integer count as the shared resource.
- The increment **method is synchronized,**
 - **Ensuring that only one thread can execute it at a time, preventing concurrent modifications.**

Mutlithreaded Counter

```
public class CounterSyncMethod {  
    public static void main(String[] args) {  
        Counter cnt = new Counter(); // Shared resource  
        Thread t1 = new Thread(() -> {  
            for (int i = 0; i < 1000; i++) { cnt.inc(); } });  
        Thread t2 = new Thread(() -> {  
            for (int i = 0; i < 1000; i++) { cnt.inc(); } });  
  
        t1.start(); t2.start(); // Start both threads  
        try { t1.join(); t2.join();  
        } catch (InterruptedException e) { e.printStackTrace(); }  
  
        System.out.println("Counter: " + cnt.get());  
    }  
}
```

Mutlithreaded Counter: Sync Block

```
// Java Program to demonstrate synchronization block in Java
class Counter {
    private int c = 0; // Shared variable
    // Method with synchronization block
    public void inc() {
        synchronized(this) { // Synchronize only this block
            c++;
        }
    }
    public int get() { return c; }
}
```

Another Example: Ticket Booking

```
// thread synchronization for Ticket Booking System
class TicketBooking {
    private int availableTickets = 10; // Shared resource (available tickets)
    // Synchronized method for booking tickets
    public synchronized void bookTicket(int tickets) {
        if (availableTickets >= tickets) {
            availableTickets -= tickets;
            System.out.println("Booked " + tickets + " tickets, Remaining tickets: " +
availableTickets);
        } else {
            System.out.println("Not enough tickets available to book " + tickets);
        }
    }
    public int getAvailableTickets() { return availableTickets; }
}
```

Another Example: Bank Balance

```
class BankAccount {  
    // Java Program to demonstrate Process Synchronization  
    private int balance    = 1000; // Shared resource (bank balance)
```

```
    // Synchronized method for deposit operation
```

```
    public synchronized void deposit(int amount) {  
        balance += amount;  
        System.out.println("Deposited: " + amount + ", Balance: " + balance);  
    }
```

```
    // Synchronized method for withdrawal operation
```

```
    public synchronized void withdraw(int amount) {  
        if (balance >= amount) { balance -= amount;  
            System.out.println("Withdrawn: " + amount + ", Balance: " + balance);  
        }  
        else { System.out.println("Insufficient bal to withdraw: " + amount); }  
    }
```

```
    public int getBalance() { return balance; }
```

```
}
```

Synchronized Method using Anonymous Class

```
import java.io.*;
class Test {
    synchronized void test_func(int n) {
        // synchronized method
        for (int i = 1; i <= 3; i++) {
            System.out.println(n + i);
            try { Thread.sleep(100); }
            catch (Exception e) { System.out.println(e); }
        }
    }
}
```

Driver Code

```
public class SyncPrintAsync {  
    public static void main(String args[]) {  
        // only one object  
        final Test O = new Test();  
  
        Thread a = new Thread() {  
            public void run() { O.test_func(15); }  
        };  
  
        Thread b = new Thread() {  
            public void run() { O.test_func(30); }  
        };  
  
        a.start();  
        b.start();  
    }  
}
```


Another Sync Example

```
import java.io.*;
// A Class used to send a message
class Sender {
    public void send(String msg) {
        System.out.println("Sending " + msg);
        try { Thread.sleep(100); }
        catch (Exception e) {
            System.out.println("Thread interrupted.");
        }
        System.out.println(msg + "Sent");
    }
}
```

Another Sync Example: Contd

// Class for sending a message using Threads

```
class ThreadedSend extends Thread {
```

```
    private String msg;
```

```
    Sender sender;
```

// Receives a message object and a string message to be sent

```
    ThreadedSend(String m, Sender obj) {
```

```
        msg = m;    sender = obj;
```

```
    }
```

```
    public void run() { // Only one thread can send a msg at a time.
```

```
        synchronized (sender) { // Synchronizing the send object
```

```
            sender.send(msg);
```

```
        }
```

```
    }
```

```
}
```

Another Sync Example: Contd

// Driver class

```
class SyncMessg {  
    public static void main(String args[]) {  
        Sender send = new Sender();  
        ThreadedSend S1 = new ThreadedSend("Hi ", send);  
        ThreadedSend S2 = new ThreadedSend("Bye ", send);  
        // Start two threads of ThreadedSend type  
        S1.start();  S2.start();  
        // Wait for threads to end  
        try {  S1.join();  S2.join();  }  
        catch (Exception e) {  System.out.println("Interrupted");  }  
    }  
}
```